

BMI / CS 771: Homework Assignment 3 Report

Qinxinghao Chen

Handan Hu

Chongwei Liu

Bohan Wen

November 15, 2025

1 Team Contributions

- **Qinxinghao Chen:** Responsible for the model implementation and training on VOC dataset.
- **Handan Hu:** Responsible for training and evaluating on COCO dataset.
- **Chongwei Liu:** Responsible for the model inference.
- **Bohan Wen:** Responsible for final model training data processing and report writing.

2 Implementation Details and Challenges

During the implementation of the FCOS model, we encountered several key technical challenges, which we resolved through a series of strategic adjustments and debugging:

1. **Training Stability:** In our initial training attempts, we observed gradient explosion, which caused the loss to become NaN. We successfully suppressed the NaN values by adopting a **lower initial learning rate (0.01)**, ensuring a stable training launch.
2. **Center-ness Loss Convergence:** We found that the *ctr_loss* converged much more slowly than other losses (see Figure 4). To address this, we adjusted the loss function by doubling its weight ($final_loss = cls_loss + reg_loss + 2 \times ctr_loss$), thereby forcing the model to pay more attention to the quality of predicted box centers.
3. **Inference Top-k Logic:** In the *inference* function, we discovered that one cannot simply flatten all predictions and take the global top-k. Instead, we modified the logic to select candidates **per-class**, which ensured that rare categories could be correctly detected.
4. **Coordinate Format Confusion:** We experienced confusion between the (x, y) and (y, x) formats when handling points, which led to significant debugging difficulties. We spent extra time unifying the coordinate system to resolve this issue.
5. **Weight Initialization:** To aid the convergence of *ctr_loss*, we applied specific initializations to the bias terms in the `RegressionHead` (e.g., initializing `bbox_ctrness.bias` to 0.0).

3 Model Inference

We first validated the correctness of our inference code using the provided pre-trained model (`voc_res18.pth.tar`).

- **mAP Score:** Our implementation achieved an mAP@0.5 of **60.1%**, which is consistent with the assignment's benchmark of 60.9%.
- **Inference Efficiency:** With a batch size of 32, the average inference time was approximately 1.16 seconds per batch, and the total evaluation time was 191.19 seconds.
- **Sample Detection Results:** As shown in Figures 1 and 2, the pre-trained model demonstrated strong detection capabilities, accurately identifying and localizing objects of various classes, such as horses, people, and cars.



Figure 1: Pre-trained model detection result.



Figure 2: Pre-trained model detection result.

4 Model Training

After validating our inference code, we trained our own FCOS model on the VOC 2007 dataset using a ResNet-18 backbone.

4.1 VOC Dataset

4.1.1 Key Parameters & Training Efficiency

Baseline (as reported).

- **Configuration:** ResNet-18, 10 epochs, Learning Rate: 0.01.
- **Efficiency:** We completed the training using a **single T4 GPU in 25 minutes**. This is “well under” the 30-minute reference standard mentioned in the assignment, proving that our `compute_loss` implementation is efficient.

Extended Training on Top of the Baseline. On top of the above baseline setting, we further trained with a lower learning rate and more epochs while keeping all other hyperparameters unchanged unless noted:

- **Configuration:** ResNet-18, 50 epochs, Learning Rate: 8×10^{-4} (0.0008).
- **Results:**

```
Running per image evaluation...
Evaluate annotation type *bbox*
DONE (t=5.87s).
Accumulating evaluation results...
DONE (t=1.45s).
Average Precision (AP) @[ IoU=0.50:0.95 | area=   all | maxDets=100 ] = 0.313
Average Precision (AP) @[ IoU=0.50      | area=   all | maxDets=100 ] = 0.574
Average Precision (AP) @[ IoU=0.75      | area=   all | maxDets=100 ] = 0.307
Average Precision (AP) @[ IoU=0.50:0.95 | area= small  | maxDets=100 ] = 0.044
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.157
Average Precision (AP) @[ IoU=0.50:0.95 | area= large  | maxDets=100 ] = 0.406
Average Recall    (AR) @[ IoU=0.50:0.95 | area=   all | maxDets= 1 ] = 0.309
Average Recall    (AR) @[ IoU=0.50:0.95 | area=   all | maxDets=10 ] = 0.424
Average Recall    (AR) @[ IoU=0.50:0.95 | area=   all | maxDets=100 ] = 0.431
Average Recall    (AR) @[ IoU=0.50:0.95 | area= small  | maxDets=100 ] = 0.082
Average Recall    (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.270
Average Recall    (AR) @[ IoU=0.50:0.95 | area= large  | maxDets=100 ] = 0.534
All done! Total time: 124.10 sec
```

Figure 3: Test result on extended model.

4.1.2 Training Curves

As shown in Figure 3, the total loss steadily decreased while the mAP on the test set steadily increased. As shown in Figure 4, the component losses (cls_loss, reg_loss) decreased significantly, while the ctr_loss decreased more slowly, validating our decision to increase its weight.




Figure 4: Tensorboard result

4.1.3 Final mAP Score (Testing mAP)

Our final model achieved an mAP@0.5 of **57.4%** on the test set.

4.1.4 Sample Detection Results

As shown in Figures 5 and 6, our model is capable of accurately detecting and localizing multiple objects in an image, such as people, dogs, sofas, and dining tables.




Figure 5: Our final model's detection result.



Figure 6: Our final model's detection result.

4.2 COCO Dataset

4.2.1 Implementation Description

- Dataset (`libs/dataset.py`):

- Added `COCODetection` class
- Implemented `get_cls_names()` returning 80 COCO class names
- Modified `build_dataset()` to handle COCO’s directory structure and annotation format.
- **Configuration** (`configs/coco_fcos.yaml`):
 - Updated paths of images and annotations
 - Used ResNet-18 as the lightweight backbone
 - Set training to 4 epochs with SGD optimizer and learning rate 0.0005
 - Set num_classes: 80

4.2.2 Experimental Procedures

1. **Data Preparation:** Downloaded COCO 2017 using `download_coco.sh` and extracted files into `data/coco/`, including training, validation, test images and annotations.
2. **Training Process:** Train the model using `configs/coco_fcos.yaml`. Training curves show convergence: classification and regression losses decrease from 0.7 to 0.3–0.4, while center-ness loss stabilizes around 0.6–0.7.




Figure 7: Training loss curves on COCO

4.2.3 Performance

- **Accuracy:** $\text{mAP}@0.5:0.95 = 0.161$, $\text{mAP}@0.5 = 0.291$, $\text{mAP}@0.75 = 0.159$. As shown in Figure 8.
- **Efficiency:**
 - Training time: 3.01 hours for 4 epochs (147,854 steps)
 - Evaluation time: 272.31 seconds on the validation set
- **Analysis:** The model achieves reasonable performance given the limited training (4 epochs). $\text{mAP}@0.5$ of 29.1% demonstrates feasibility on COCO with minimal training. With more epochs and hyperparameter tuning, performance would likely improve.

```

Running per image evaluation...
Evaluate annotation type *bbox*
DONE (t=24.88s).
Accumulating evaluation results...
DONE (t=4.79s).
Average Precision (AP) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.161
Average Precision (AP) @[ IoU=0.50 | area= all | maxDets=100 ] = 0.291
Average Precision (AP) @[ IoU=0.75 | area= all | maxDets=100 ] = 0.159
Average Precision (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.050
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.174
Average Precision (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.260
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 1 ] = 0.171
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 10 ] = 0.271
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.285
Average Recall (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.072
Average Recall (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.319
Average Recall (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.462
All done! Total time: 272.31 sec

```

Figure 8: Evaluation result on COCO

5 Conclusion and Future Work

Summary. We implemented an end-to-end FCOS detector, validated inference correctness against the provided `voc_res18.pth.tar` and reproduced a strong baseline ($mAP@0.5 = 60.1\%$). On VOC, the baseline training (ResNet-18, 10 epochs, $LR = 0.01$) finished in **25 minutes** on a single T4, meeting the efficiency target. Building on this, an extended schedule (**50 epochs**, $LR = 8 \times 10^{-4}$) delivered stable convergence; our final model achieved **$mAP@0.5 = 57.4\%$** on the test set. On COCO (4 epochs, ResNet-18), we obtained **$mAP@[.5:.95] = 0.161$** , **$mAP@0.5 = 0.291$** , showing reasonable performance under a short schedule.

Key takeaways. (1) Reweighting the centerness loss (factor $2\times$) accelerates its convergence and improves localization quality; (2) Per-class candidate selection in inference avoids head-class domination and benefits rare categories; (3) Unifying (x, y) vs. (y, x) conventions eliminated several failure cases.